

Dinghy: Horizontally Scaling Raft Clusters

Drew Ripberger

January 26, 2019

Abstract

With the introduction of Raft, distributed consensus has become more widely available in the use of cluster design. Different applications of distributed systems often require differing configurations, however, as these clusters begin to increase in the number of participating nodes, they become increasingly less efficient at coming to consensus. Current methods of scaling distributed systems generally implement some variation of data sharding, batch processing, or message coalescing. We investigate however, utilizing dynamically set timeouts and heartbeat intervals to increase network throughput, and propose methods to more effectively handle a growing number of nodes in a cluster, while analyzing their ability to increase the horizontal scalability of a Raft cluster.

1 Introduction

The Raft consensus algorithm originated to simplify the preexisting Paxos algorithm, while at the same time, solving the same core problem [17] with a similar efficiency. For years, Paxos had dominated distributed consensus. At its core it defined a way in which a system could come to agreement on a given state [13]. Though, Paxos can be incredibly hard to comprehend. Many papers have been published in an attempt to offer a clearer explanation as to how Paxos functions [14, 15], but it continues to be a difficult system to implement at a practical

level.

Ultimately, these algorithms define a method for a system to agree on a state [9]. They work to build a fault tolerant approach to distributed systems, the *replicated state machine* [18]. In this context, a group of machines replicate a single state across themselves to create a fault tolerant system, that can handle the failure of $n/2 - 1$ nodes. The essential goal of consensus in terms of the *replicated state machine* is to reach a *univalent* state, from any *multivalent* state. Such algorithms specifically order state changes, to ensure that when applied, that all result in the same state [12, 17]. Raft also works to correct, and right, any nodes in a cluster in contradicting states. It does this via counting **election terms**, demonstrated as such:

We define two nodes in a cluster, with two corresponding state machines, M and N . We also define a function, $T(S)$, of some arbitrary state machine S , that is its current **term**, where T_c is the correct **term** in the cluster.

$$T_c = \begin{cases} M, & T(M) > T(N) \\ N, & T(M) < T(N) \\ \emptyset, & T(M) = T(N) \end{cases}$$

In order to keep some sense of order in the cluster, Raft keeps track of the number of leader elections that have occurred with the **election term**. This is a system wide tally that is used to determine when a node may be behind or have conflicting information in its log. Many of these comparisons have to be made from node to node through heartbeat messages,

that also act to check for leader liveness [17]. But, as one might imagine, as you try to include more nodes in a cluster, the number of checks that have to be propagated to ensure an effective system drastically increases.

2 Scaling Distributed Systems

In practical applications, clusters of varying sizes are required. In some cases many nodes will be used, each replicating a small piece of data many times over. While in others, few nodes will be used and larger chunks of data are replicated. Though in implementation, there are drawbacks to having a cluster with many nodes. As you continue to increase node number, various factors can lower a network's required time to reach consensus. Raft solves the consensus problem algorithmically, but let's, for example, take a look at a real world example where scaling comes into play.

We can imagine a large party of friends trying to decide where they want to go to eat for dinner. In this scenario in order to make an effective, and satisfying decision as to where the group should dine, each member of the group must be consulted. So the time in which it takes the entire group to come to an agreement increases as the size of the group increases.

In this large party, many more people will have to be consulted, and more options will have to be weighed before a choice is made. Though, compare this to just a few friends, who would be able to reach mutual agreement much faster, as they have less to consider, and fewer people that need to be taken into account before reaching a decision.

This principle is clearly demonstrated in distributed systems [16]. Adding more nodes to a cluster makes it more difficult for the network to handle faults and replicate its state. Demonstrated in Raft, the more *followers* in a system, the more heartbeats that have to be sent out to the *leader*, processed, responded to, and then confirmed [17]. This will be demonstrated realistically later.

Now we begin to look at some of the possible

methods to scale distributed systems, each with varying degree of complexity and success.

2.1 Non-Voting Servers

One common way of increasing the capabilities of a large Raft cluster, is to change the voting status of additional nodes. Usually when we discuss creating larger clusters, we mean creating a cluster of servers, each with full voting privileges. But every one of these members all having full voting power, requires extra considerations in the cluster. When you are a *full* node, you must participate in log replication, as well as be consulted when changes need to be made to the replicated log. So a common solution to this requirement for being a *full* voting server, is making some of the members *non-voting* servers. This allows for the node to participate in the cluster, gaining the benefits of the log replication, but not bogging down the cluster by increasing the quorum size [2]. Though it should be noted that this solution does not in any way solve the presented problem. Demoting nodes to a non-voting configuration, does not assist with the fault tolerance of the cluster. Suppose you spin up a cluster with 5 full Raft servers, and 100 non-voting servers. If even just 3 of the full nodes fail, the system can no longer make progress, even though we have allocated resources for running a total of 105 servers because, the majority of the voting nodes have failed.

2.2 Sharding

There is also a way of scaling writes and reads. Some databases divide up the data they are trying to replicate using a process called *sharding*. A great example of this process can be found in Google's Spanner database [10]. *Sharding* works as a type of distributed load balances to balance the store of data across many consensus groups.

2.3 Batching

Other databases work to alter the process of sending or receiving inter-node messages. The Calvin database works to group messages, and

specially schedule them in order to maintain consistency, and scalability [19].

3 Scaling Raft With Dynamic Timeouts

Like all other algorithms of its class, Raft has similar problems scaling. Though it has a distinctly unique set of problems. It should be noted that we can consider Raft generally *synchronous*. That is, all nodes must reach a *univalent* state, before handling the replication of the next update to the log. Raft’s use of a *leader*, requires all changes to the network state be processed through a single node, thereby essentially creating a synchronous cluster where in which *followers* have very little power to influence the cluster’s state when there is a healthy *leader*. So when we discuss scaling in terms of a leader based methods we have to examine different properties of the algorithm.


Dinghy works as an embedded algorithm within Raft. It runs side by side the core Raft *follower* routines, in order to actively maintain historical statistics on a node. In its most base form Dinghy attempts to get a sense of the environment that a given server is running in and adjust some of its previously set, static parameters, and adjust them to a more appropriate level for the situation. The Lifeguard system works in a similar fashion, but with differing goals, and approaches, using a step function to help stop false positives, whereas Dinghy looks at historical averages [11].

4 Base Performance

Of course before we look into some potentially beneficial alterations we could make to Raft, we must examine the base algorithms ability to scale.

$$L(n) = 1 - \sum_{i=0}^{n/2-1} \binom{n}{i} (0.75)^i (0.25)^{n-i}$$

$$\{(n, p) : n \in \{3, 5, 7, \dots, 51\}, p = L(n)\}$$



5 Testing Consensus

The big draw of Raft is its practicality. It has allowed many the opportunity to implement distributed consensus, as it was designed to be understandable. So with this in mind, we maintain a similar dogma when it comes to testing the algorithms scalability. We wanted to design a test that would stress test the networks throughput given a certain size, given that Raft is overwhelmingly used in creating fault tolerant databases [1, 4, 5, 7].

With these considerations, we first decided the most common use case for this method of distributed consensus. The official Raft website contains a useful list of implementations [8]. Given its capabilities, Raft is overwhelmingly used for database replication. We elected to use a popular open source library written in Go, `hashicorp/raft` [6]. With this, we used an integer state in our cluster’s state machine. We define the function $S(R)$, that is the current state of some arbitrary Raft node, R .

Our testing tracks the number of successful **Apply** operations we are able to make to the replicated state in some time. The time allotted for a single test is denoted as T_f . After every state update, a barrier is called in order to ensure all of the nodes are consistent. After $T_f/4$ the leader of the test cluster is killed. That is to say, it is non-gracefully shutdown, will stop

returning messages, and will not return to be a functioning member of the cluster at any time during the test. The replicated state machine is being incremented every update, such that, after T_f , given a set of Raft nodes C ,

$$\{r_1, r_2 \in C : S(r_1) = S(r_2)\}$$

and the final number of applies, or commits for our test was the value of the final integer state of some arbitrary node in the cluster.

In this specific instance we elected for a shorter testing time of 5 seconds, then of course by our procedure, the leader would be killed after 1.25 seconds [3]. We use this specific procedure in the testing of both our proposed algorithm, Dinghy, and the benchmarks of pure Raft.

References

- [1] Cockroachdb - the open source, cloud-native sql database. <https://github.com/cockroachdb/cockroach>. Accessed: 2018-11-15.
- [2] Consul enterprise enhanced read scalability. <https://www.consul.io/docs/enterprise/read-scale/index.html>. Accessed: 2019-1-24.
- [3] Dinghy: Horizontally scaling raft clusters. <https://github.com/drewrip/dinghy>. Accessed: 2019-1-24.
- [4] Distributed reliable key-value store for the most critical data of a distributed system. <https://github.com/etcd-io/etcd>. Accessed: 2018-11-15.
- [5] Distributed transactional key-value database, originally created to complement tidb. <https://github.com/tikv/tikv>. Accessed: 2018-11-15.
- [6] Golang implementation of the raft consensus protocol. <https://github.com/hashicorp/raft>. Accessed: 2018-11-15.
- [7] The open-source database for the realtime web. <https://github.com/rethinkdb/rethinkdb>. Accessed: 2018-11-15.
- [8] Raft consensus algorithm. <https://raft.github.io>. Accessed: 2018-11-15.
- [9] BRAND, D., AND ZAFIROPULO, P. On communicating finite-state machines. *J. ACM* 30, 2 (Apr. 1983), 323–342.
- [10] CORBETT, J. C., DEAN, J., EPSTEIN, M., FIKES, A., FROST, C., FURMAN, J., GHEMAWAT, S., GUBAREV, A., HEISER, C., HOCHSCHILD, P., HSIEH, W., KANTHAK, S., KOGAN, E., LI, H., LLOYD, A., MELNIK, S., MWAURA, D., NAGLE, D., QUINLAN, S., RAO, R., ROLIG, L., SAITO, Y., SZYMANIAK, M., TAYLOR, C., WANG, R., AND WOODFORD, D. Spanner: Google’s globally-distributed database. In *10th USENIX Symposium on Operating Systems Design and Implementation (OSDI 12)* (Hollywood, CA, 2012), USENIX Association, pp. 261–264.
- [11] DADGAR, A., PHILLIPS, J., AND CURREY, J. Lifeguard : Swim-ing with situational awareness. *CoRR abs/1707.00788* (2017).
- [12] LAMPORT, L. Time, clocks, and the ordering of events in a distributed system. *Commun. ACM* 21, 7 (July 1978), 558–565.
- [13] LAMPORT, L. The part-time parliament. *ACM Trans. Comput. Syst.* 16, 2 (May 1998), 133–169.
- [14] LAMPORT, L. Paxos made simple. *ACM SIGACT News (Distributed Computing Column)* 32, 4 (Whole Number 121, December 2001) (December 2001), 51–58.
- [15] MAZIERES, D. Paxos made practical. *Unpublished manuscript, Jan* (2007).
- [16] NEUMAN, C. Scale in distributed systems. *Readings in Distributed Computing Systems* (1994).
- [17] ONGARO, D., AND OUSTERHOUT, J. In search of an understandable consensus algorithm. In *Proceedings of the 2014 USENIX Conference on USENIX Annual Technical Conference* (Berkeley, CA, USA, 2014),

USENIX ATC'14, USENIX Association, pp. 305–320.

- [18] SCHNEIDER, F. B. Implementing fault-tolerant services using the state machine approach: A tutorial. *ACM Comput. Surv.* 22, 4 (Dec. 1990), 299–319.
- [19] THOMSON, A., DIAMOND, T., WENG, S.-C., REN, K., SHAO, P., AND ABADI, D. J. Calvin: Fast distributed transactions for partitioned database systems. In *Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data* (New York, NY, USA, 2012), SIGMOD '12, ACM, pp. 1–12.