Scaling Raft For Varying Cluster Sizes

Drew Ripberger

October 28, 2018

Abstract

Placeholder for the abstract

1 Introduction

The Raft consensus algorithm originated to simplify the preexisting Paxos algorithm, while solving the same core problem [4] with the same efficiency. For years, Paxos had dominated distributed consensus. At its core it defined a way in which a system could come to agreement on a given state [1]. Though, Paxos can be incredibly hard to comprehend. Many papers have been published in an attempt to offer a clearer explanation as to how Paxos functions [2, 3], but it continues to be a difficult system to implement at a practical level.

References

[1] Lamport, L. The part-time parliament. *ACM Trans. Comput. Syst. 16*, 2 (May 1998), 133–169.

[2] Lamport, L. Paxos made simple. *ACM SIGACT News (Distributed Computing Column) 32, 4 (Whole Number 121, December 2001)* (December 2001), 51–58.

[3] Mazieres, D. Paxos made practical. *Unpublished manuscript, Jan* (2007).

[4] Ongaro, D., and Ousterhout, J. In search of an understandable consensus algorithm. In *Proceedings of the 2014 USENIX Conference on USENIX Annual Technical Conference* (Berkeley, CA, USA, 2014), USENIX ATC'14, USENIX Association, pp. 305–320.